

Stats 545: Homework 5

Due before class, Thursday Nov 8.
All plots should have labelled axes and titles.

Important: Rcode, tables and figures should be part of a single .pdf or .html files from R Markdown and knitr. See the class reading lists for a short tutorial. Any derivations can also be in Markdown, in Latex or neatly written on paper which you can give to me.

1 Problem 1: Exponential family distributions [20]

This continues from parts (a) and (b) Problem 1 from Homework 4 on exponential family distributions.

1. Consider a random variable x that can take D values and that is distributed according to the discrete distribution with parameters $\vec{\pi}$. We will write this as $p(x|\vec{\pi})$, with $p(x = c|\vec{\pi}) = \pi_c$ for $c \in \{1, \dots, D\}$.
 - (c) What is the form of the conjugate prior on the parameters of $p(x|\pi)$? You only need to write it upto a multiplicative constant (i.e. you don't have to write the normalization constant). What is its feature vector? [3 pts]
 - (d) If you call the natural parameters of this distribution $\vec{a} = (a_1, \dots)$ (part (c) will give the dimensionality of \vec{a}), what are the parameters of the posterior distribution given a set of observations $X = (x_1, \dots, x_N)$? (The point here is that in general the posterior distribution can be very complicated, even for simple priors (so that we need methods like MCMC). However for conjugate priors, the posterior lies in the same family as the prior, it just has different parameters) [2 pts]
2. Let x be Poisson distributed with mean λ . Repeat parts (a), (b). [5 pts]
3. Let x be a 1-dimensional Gaussian with mean μ and variance σ^2 . Repeat parts (a), (b) (Note: both μ and σ^2 are parameters). [5 pts]
4. Let x follow a geometric distribution with success probability p : ($\Pr(X = k) = (1 - p)^k p$ for $k = 0, 1, 2, \dots$). Repeat parts (a), (b). [5 pts]

2 Problem 2: Gradient descent for blind source separation [100 pts]

Let \mathbf{Y} be a $3 \times T$ matrix consisting of three independent audio recordings. Each column $\mathbf{y}_t = (y_{1t}, y_{2t}, y_{3t})^\top$ consists of the intensities of each source at time t . Even though each audio signal is a time-series, we will model the y_{ij} 's as i.i.d. draws from the hyperbolic secant distribution $p(y) = \frac{1}{\pi \cosh(y)}$.

1. Plot $p(y)$. Superpose a Gaussian with the same mean and variance. [3 pts]

Instead of observing \mathbf{Y} , we observe $\mathbf{X} = \mathbf{A} \cdot \mathbf{Y}$, where \mathbf{A} is a 3×3 'mixing'-matrix. Given \mathbf{X} , we want to recover *both* \mathbf{A} and \mathbf{Y} . Note this problem is *ill-posed* i.e. there are an infinite collection of valid (\mathbf{A}, \mathbf{Y}) pairs. However, our model of \mathbf{Y} (in particular, the source-independence and non-Gaussianity assumptions) will allow us to recover a sensible solution.

Assume, \mathbf{A} is invertible, and define $\mathbf{W} = \mathbf{A}^{-1}$. Write \mathbf{w}^s for row s of \mathbf{W} .

2. Show that $\log p(\mathbf{X}|\mathbf{W}) = T \log |\mathbf{W}| + \sum_{t=1}^T \sum_{s=1}^3 \log p(\mathbf{w}^s \mathbf{x}_t)$. (Hint: see e.g. the last paragraph here: http://sccn.ucsd.edu/wiki/Random_Variables_and_Probability_Density_Functions). [5 pts]

We will find the maximum likelihood (ML) estimate of \mathbf{W} by gradient ascent.

3. Show that $\frac{\partial \log p(\mathbf{X}|\mathbf{W})}{\partial w_{ij}} = T a_{ji} + \sum_{t=1}^T \frac{\partial \log p(y_{it})}{\partial y_{it}} x_{jt}$ (Hint: look at Wikipedia for the derivative of a determinant). [5 pts]

It follows that

$$\frac{1}{T} \frac{\partial \log p(\mathbf{X}|\mathbf{W})}{\partial \mathbf{W}} = \mathbf{A}^\top + \frac{1}{T} \sum_{t=1}^T \frac{\partial \log p(\mathbf{y}_t)}{\partial \mathbf{y}_t} \mathbf{x}_t^\top$$

For our choice of $p(y)$, we have $\frac{\partial \log p(\mathbf{y}_t)}{\partial \mathbf{y}_t} = -\tanh(\mathbf{y}_t) := [-\tanh(y_{1t}), -\tanh(y_{2t}), -\tanh(y_{3t})]^\top$.

The summation over T on the RHS can be written as $-\tanh(\mathbf{Y}) \cdot \mathbf{X}^\top$: this is just one line of R code.

4. Gradient ascent iteratively tries to reach a local maximum of the likelihood according to the rule

$$\mathbf{W}_{new} = \mathbf{W}_{old} + \eta \left(\frac{1}{T} \frac{\partial \log p(\mathbf{X}|\mathbf{W})}{\partial \mathbf{W}} \bigg|_{\mathbf{W}_{old}} \right)$$

The last term is the gradient evaluated at \mathbf{W}_{old} . Plug in the expression for the gradient to write the update rule. [5 pts]

5. Write a function that takes a matrix \mathbf{X} as input, and performs gradient ascent to get an ML-estimate of \mathbf{W} . Your function should have an initialization rule and a stopping criteria (which you will describe later). You will also need to set the learning rate η . You can get better results by using a smaller η for later iterations. [15 pts]

On the course webpage are three wav files, `mike1.wav`, `mike2.wav` and `mike3.wav`. Download them, and load them into R using `load.wave` from the `audio` package. Hopefully, you'll get three 490000-length vectors, which you can store into a matrix \mathbf{X} . Directly running your function on this probably won't work (but you can try).

6. Write the 3×3 covariance matrix of \mathbf{X} . Also, for each pair of recordings, plot a scatterplot of values (i.e. for rows (1,2), (2,3) and (1,3)). [10 pts]
7. If `Cov` is the covariance matrix, calculate its square root using the `sqrtn` function from the `expm` package. Call this `sqCov`. Define a new 'whitened' matrix $\mathbf{X}_{white} = \mathbf{sqCov}^{-1} \cdot \mathbf{X}$. This should have covariance equal to the identity matrix. Verify this. [10 pts]
8. Run your gradient ascent algorithm on the whitened data. How did you initialize \mathbf{W} ? Describe how you set the learning rate η as well as your stopping criterion (I set η to about 1). How many iterations did you need? Plot the evolution of the log-likelihood. What is the returned estimate $\hat{\mathbf{W}}$? Your actual demixing matrix for the original data \mathbf{X} is this value multiplied by \mathbf{sqCov}^{-1} . Print this out. If you had to restart with different initializations, describe this too. [30pts]
9. Plot histograms of the raw data, the whitened data and the recovered data (a total of 9 histograms) [9pts]
10. Your estimate of the latent sources is given by $\hat{\mathbf{Y}} = \hat{\mathbf{W}} \cdot \mathbf{sqCov}^{-1} \cdot \mathbf{X}$. Plot a histogram of the marginals for each source. What is their covariance? Plot the three pairwise scatterplots. [5pts]

You can verify your answer by listening to the recovered sources. You can save these using the `save.wave` command. You should normalize these first (I divided by twice the maximum value for each source).

11. Is the maximum likelihood estimate of \mathbf{W} you obtained unique? If so explain why, if not, explain some sources of degeneracy. [2 pts]